

Operating Systems Homework 1 Report

Andres Ponce, 0616110

2020-10-13

1. What is a kernel? What are the differences between *mainline*, *stable*, and *longterm*? What is a kernel panic

The **kernel** of the operating system refers to the program that is always running on the computer. The kernel might include software for CPU scheduling, file system management, etc... which are integral to the normal operation of the system. There are different ways to release an operating system kernel.

The **mainline** kernel is the kernel that is currently being worked on and developed. Since there is constant development work occurring on the mainline kernel, there are constant releases of this kernel.

Once the mainline kernel has been released and iterated on, it moves to be a **stable** kernel. The stable kernel receives less updates than the mainline kernel, those usually being more significant bug fixes.

The final step is becoming a **longterm** kernel. This kernel will be mostly be used for bugfixes for older versions of the operating system.

A big purpose of these releases and constant updates is to minimize the occurrence of **kernel panics**. These are errors which might have serious consequences on the operation of the kernel. The causes for a kernel panic might involve unrecoverable errors in memory, drivers, or other kernel component.

2. What are the differences between *building*, *debugging*, and *profiling*?

Building the kernel refers to compiling the source operating system source code. In the second step of the homework, we download the kernel using the `wget` command. Then we further download additional required dependencies and compile the kernel source code using the `sudo make -j$(nproc)` command. This turns all the code in the corresponding Linux version to a bootable format.

Once we have built the kernel, we have to **debug** it. Like in any program, the operating system is bound to have some flaws or unintended behavior, or *bugs*. However, since we are compiling a pretty basic layer of software on which all our other programs run, it is quite different. To debug a kernel, we require one machine to make our changes and then we send our changes to another machine.

While we first debug our program and then build it, we still have to measure its performance, which we call **profiling**. This process involves tracing the performance or measuring the number of system calls to identify potential performance bottlenecks. Programs such as Valgrind can trace the amount of system memory used, and other such programs can help us measure how well our programs perform.

3 What are GCC, GDB, and KGDB, and what are they used for?

GCC refers to the **GNU Compiler Collection**, a collection of frontends for various languages. However, if in bash we just type `gcc` we will get the C compiler to run. We use this specific compiler to compile the kernel, since much of the source code is written in C.

GDB refers to the GNU project's debugger. Essentially, we can run the program using certain precautions. For example, GDB can start our program and place breakpoints within it. We can also check the values

of variables at that point, along with performing small changes to debug. For this project, we perform some debugging on the kernel using GDB and then send the changes to the other virtual machine. In this way, we go around the issue of debugging the operating system we need to debug the operating system.

KGDB is a debugging tool used to debug the kernel. For example, in the assignment, we set the `CONFIG_FRAME_POINTER` setting inserts code directly in the executable. This code then can save the state of the registers during execution. We can then use this code with GDB during the debugging process.

4. What are the `/usr/`, `/boot/`, `/home/`, `/boot/grub` folders for?

These folders are all directories under the Linux filesystem. They all have different purposes. For example the `/usr/` directory contains user utilities that are shared among all users of a system. The `ls` program which lists the files in a directory is located in `/usr/bin`, for example.

The `/boot/` and `/boot/grub` folders concern the programs and procedures needed at boot time. The regular `/boot/` folder contains code for UEFI (depending on the system) or BIOS, the code for `initramfs` which decompresses the kernel during the boot process, as well as the actual Linux image file. During the boot process, we use the bootloader to load the initial filesystem in main memory, which is in charge of loading the main kernel.

Inside the `/boot/` directory there also resides the code for `GRUB`, which stands for the **GR**and **U**nified **B**ootloader. The bootloader is a program responsible for loading the main operating system kernel. GRUB can detect the operating system(s) present on a machine and provides a way to select which operating system to load. It also allows a restricted command line, whose commands are also defined under `/boot/grub`.

The `/home` directory is where most of the individual user's files are kept. When we spawn a new terminal, it will place us in the user's home folder, whose contents are only available to the currently signed in user.

5. What are the general steps to debug a Linux kernel?

Firstly, we need two instances of Linux to debug the kernel. We have to use one instance to run actually debug the kernel and then we send the changes to the other Linux instance. We use GDB to remotely debug the kernel.

The Target machine is where all our kernel source code is stored. We build the kernel and copy the image file from the Target machine to the Host machine using `scp`. Then upon restarting the Target machine we are brought to the `kdb` prompt before we log in to the system. At this point, from the Host machine, we can run `gdb ./vmlinuz` to start debugging the kernel image. Then, we set the target of the debug to be the Target machine using the target machine's IP address. We test this by using the

continue command on the gdb prompt and waiting for our Target machine to boot into the Ubuntu menu.

The overall debug process follows these general steps:

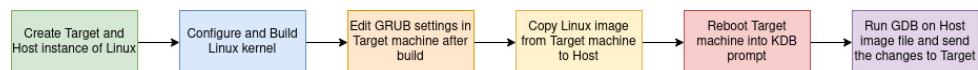


Figure 1: Debug Process

- Create two virtual machines, one for running the kernel and the other for the actual debugging.
- Edit the settings in the Target machine to support debugging
- Build the kernel.