

Chapter 3: Processes

Andrés Ponce

October 27, 2020

A **process** is a program that is currently running. One of the main functions of an operating system is to run user programs, so the way an OS manages process execution, scheduling and memory management can become quite important.

The status of the current executing program can be known with the value of the **program counter** ¹ and the values currently in the registers.

Usually, a program executable file has four distinct sections. The **text section** contains the executable code, while the **data section** contains the global variable declarations. the **heap** is where the dynamic memory is allocated, and the **stack** is where temporary variables are allocated. These variables are the ones used for function executions and are then deleted.

The data and text sections of a program are fixed, since they do not change during execution. However, the stack and heap shrink as needed, and they grow **toward** each other but never **overlap** each other. ²

The program state can fall into certain categories:

- **New:** Process is being created.
- **Running:** There are instructions being executed.
- **Waiting:** Process is waiting for some event.
- **Ready:** Process is ready to be assigned to CPU.
- **Terminated:** Process finished execution.

Process Control Block

The process is represented in the OS by using a block of memory containing program info called the **Process Control Block**. It stores a couple things:

- **Process State:** One of previously mentioned process state.
- **Program Counter:** Address of the next instruction to be executed.
- **CPU registers:** There are different types of registers that store different information. This has to be saved on interrupt to be able to resume when the interrupt is resolved.

¹ The PC stores the address of the currently executing instruction.

² Then how do we account for really big programs?

- **CPU scheduling:** Scheduling parameters.
- **Mem. Mgmt. Information:** The value of the base and limit registers, page tables, segment tables, depending on memory scheme in computer.
- **Accounting Information:** Amount of CPU and real time used, time limits, process/job numbers, etc...
- **I/O Status Information:** List of open files, I/O devices allowed to use, etc...

The set of PCB's is stored as a linked-list, with the info stored in `include/linux/sched.h`. When we want to schedule a process, we have to change the state. Then we execute until we are finished or waiting for a process, and move to the next program in the **wait queue**.

In terms of scheduling, each process gets a CPU core when it needs to and gets removed when not in need of execution, such as waiting for I/O. For CPU bound instructions though, they also don't get to execute constantly throughout its lifetime, the CPU scheduler might forcibly remove it after a specified time limit.

When changing the CPU core from execution of one process to another, most systems will save the current state of the executing program and load the state of the program to execute. This switching of process is known in the biz as **context switching**, since the state of the registers and execution point of a program when saved and loaded can be called the context.

Operations on Processes

When executing, each process on Windows and UNIX gets a unique **process identifier**,³ which uniquely identifies the process in the kernel. Since each process can spawn other children processes, the execution diagram can become an execution **tree**. In Linux, the **systemd** process always has a pid of 1. This process is the only user program run at boot time, and every process starts under the systemd node in the execution tree. The **logind** process manages the clients that are logged in to the system.

What happens when we create/run a program? First, the OS is responsible for allocating the resources the process needs to operate. Remember that a process can have another process execute inside of it called a **child process**. Since the parent has some execution and read/write permissions, as well as some data, the child process inherits some of them. However, the child's processes are to be a **subset** of the parents', and this is done to discourage the creation of unnecessary

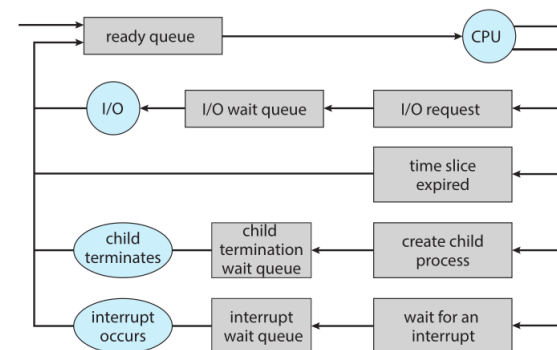


Figure 1: For an active process, the process may be not executing while it's waiting for I/O, or while it's waiting for interrupt. During execution, this program might be descheduled to allow other ready programs to execute. It's taken off the ready queue.

³ a.k.a. pid

processes. Child processes will overlay their memory footprint with that of the parents, so that they occupy some of the same address space as the parents. Also, during execution of the child process, the parent's `pid` is replaced with that of the child. The child technically has a `pid` of 0, although its true one is taken on by the parent.

How do processes terminate? When they reach the final statement to be executed, they call on the `exit()` system call. Parents can also terminate their own children, but cannot terminate other programs. Why would a parent terminate its child process? Maybe the child's available resources have been exceeded, or the child process is simply no longer needed. If the parent is executing, we also need the child process to execute as well, since the OS usually does not permit parent processes to terminate while their children are still executing.

Android Process Hierarchy

In the Android operating system, instead of terminating arbitrary processes, there is a hierarchy of importance. The motivation for this is that mobile systems need to aggressively control memory and thus need to constantly be reclaiming system resources. The categories of importance in Android includes

- **Foreground process:** This process is the one currently occupying the user's screen.
- **Visible process:** A process that is not visible on the screen but that performs a function whose result is displayed on the foreground process.⁴
- **Background process:** Process that performs an activity not apparent to the user.
- **Empty Process:** Process not associated with any application.

⁴ Are these the processes spawned by the foreground process?

Interprocess Communication in Shared Memory

Previously, we had mentioned two models for interprocess communication: shared memory and message passing. In the shared memory model, a region in memory is designated as memory that both processes can use to store info they want passed. In the message passing model, we can directly send a message between the two programs.⁵ Shared memory is usually faster than message passing because message passing can involve many system calls and be overall slower. However, the process to use shared memory is up to the processes and not the OS, so they have to both make sure they handle it correctly and are not writing to the same region simultaneously.

⁵ Is this where sockets or system links come in?

In shared memory, there is a **producer** and **consumer**. The producer makes some info that the consumer will read. For example, a compiler will produce some code that might be read by the assembler and the assembler in turn produces machine code to be executed. Thus for the producer-consumer model we also need a way to model it in the memory buffer. The two processes will use a memory buffer of certain size to communicate. This might be a **bounded** or **unbounded** buffer, in which there is no strict size limit. In case of a bounded buffer, the consumer might have to wait for new stuff from the producer and likewise the producer will have to wait if the buffer is full before putting in new information.

Interprocess Communication in Message-Passing Systems

How do we pass along information which may be used by two different systems, or two different computers. We need some **communication link** in order to send the messages. We could have a scheme in which a message-passing system calls a function like `send(P, message)` to send message `message`

Another scheme involves using **ports** or **mailboxes**. Similar to mail protocols, the mailbox and ports involve placing a message there and the message being read at some point by the other program. So in this scheme we send and receive messages from the mailbox. The owner process will start the mailbox and then the mailbox's access info might be passed to the requesting program.

Processes might either be **synchronous** or **asynchronous**. These two terms refer to whether the message has to be read by the other process before being able to write to the process again.

The messages, when they arrive at a port, will be placed in a queue. In UNIX we use the `mach_msg()` to indicate if the current port is a receiver or sender process.

In Windows, if we wish for two processes to communicate, we have to communicate with a **subsystem server**, which assigns a channel for the process to use.